

Concurrency

Week 5

Semaphores
Deadlocks

Topics

- Managing Critical Section
 - Mutex (discussed last week)
 - Semaphores
- Deadlocks
 - Definition
 - Deadlocks with Mutex
 - Deadlocks With Semaphores
- Practical Exercises

Learning Outcomes

- At the end of this lesson, you will be able to:
 1. Describe the concept of a Semaphore and how it can provide safe access to shared resources.
 2. Explain what deadlocks are and how they can be prevented
 3. Apply deadlock-free protection of shared resources using semaphores.

What is the Main Idea of the Course?

- The following four questions summarize this course.
- Questions:
 1. ***How can we run multiple tasks simultaneously?*** We discuss multitasking and multi-threading to answer this question.
 2. ***How can simultaneous tasks share resources?*** We discuss resource types and how they can be shared.
 3. ***What problems may we have when we share resources?*** We discuss race conditions and the inconsistency of data to answer this question.
 4. ***What solution do we have for these problems?*** We discuss semaphores, and mutex locks to answer this question.

Part One!

Part One!

- In this part, we will discuss semaphores and how they can be used to avoid race conditions. Try to find answers to the following questions:
 1. What operations are allowed on a semaphore?
 2. How can a semaphore guarantee mutual exclusion in a critical section?
 3. Do solutions for race conditions using semaphores avoid starvation?
 4. How are counting semaphores used in message-passing systems?
 5. How do you compare semaphore with mutex?

Semaphores

Semaphore **P()**

Semaphore: an integer combined with P and V operations

- Call P to enter a critical section
 - from Dutch:
 - $P \rightarrow$ Proberen, or Passeren, or Pakken.
- It is also known as wait operation
- **If semaphore value > 0 , it can be *decremented* by one**
 - Atomic decrement
- If semaphore value $= 0$, calling P will cause the process to wait (blocks the process/thread)

Semaphore V()

Semaphore: an integer combined with P and V operations

- Call V to exit a critical section
 - from Dutch:
 - V → Verhogen or Vrijgave
- It is also known as signal operation
- **Increments** the value of the semaphore by one
- Allows the process that is waiting on V to continue

Semaphore

Semaphore: an **integer** combined with P and V operations

- For mutual exclusions problems, semaphore values are initialized to 1. These semaphores can get only two values: 0 and 1, and therefore are called binary semaphores

```
Process P
// some code
// s is a semaphore
P(s);
    // critical section
V(s);
// remainder section
```

```
Process Q
// some code
// s is a semaphore
P(s);
    // critical section
V(s);
// remainder section
```

Example 1: Bridge

Remember the Bridge problem from the previous week?

- Assume two or more observers are counting events using a shared record named **Total**.
- Observers use **semaphores** to avoid **race conditions**
- Global *Semaphore* $S = 1$

Process Observer

if Observed():

 P(S);

 temp = Read(Total)

 temp = temp + 1

 Store(temp , Total)

 V(S);

Example 2: Travel Agency

- A sales agent who is making a reservation for a flight should
 - Decrement the value of a semaphore (P)
 - Read the reservation list
 - Update the list
 - Increment the value of the semaphore (V)
- What happens if the sales agent forgets to complete the reservation (stays inside its critical section)? Discuss.

Example 3: Printer Spool


- To add an item to the print spool a process uses the following code:
- Global *Semaphore* S initialized to 1

```
Data = Prepare()  
P(S);  
    Spool[in] = Data  
    in = in + 1  
V(s);
```

Producer-Consumer

Remember: Producers and Consumers share a buffer ...

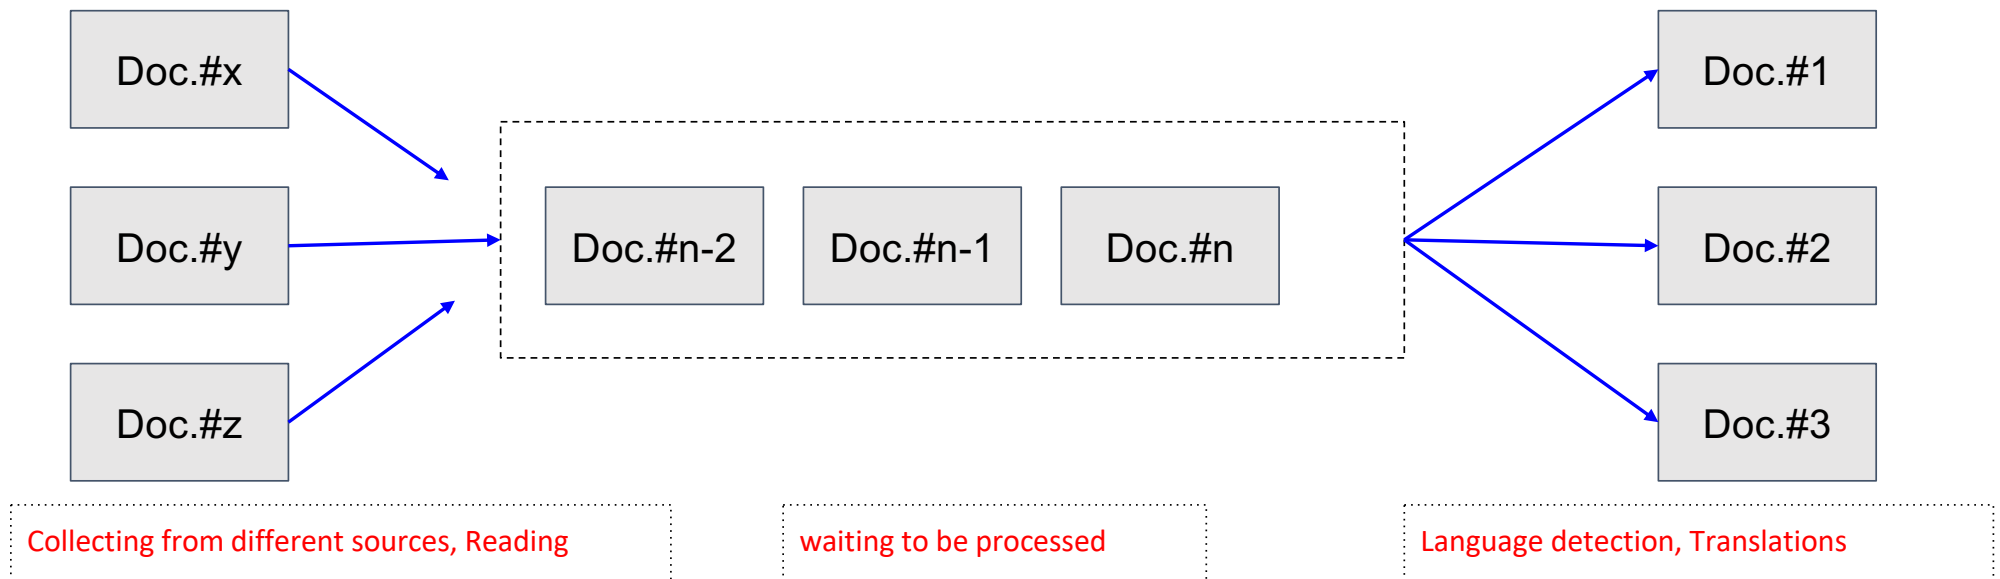
- How can Producers and Consumers inform each other about the state of the buffer?
- Consumer prefers to pick the item when it is ready, otherwise iteratively it has to check the buffer



Producer-Consumer: Example

Translating documents:

- Translator threads need to be informed when a document is ready for translation.



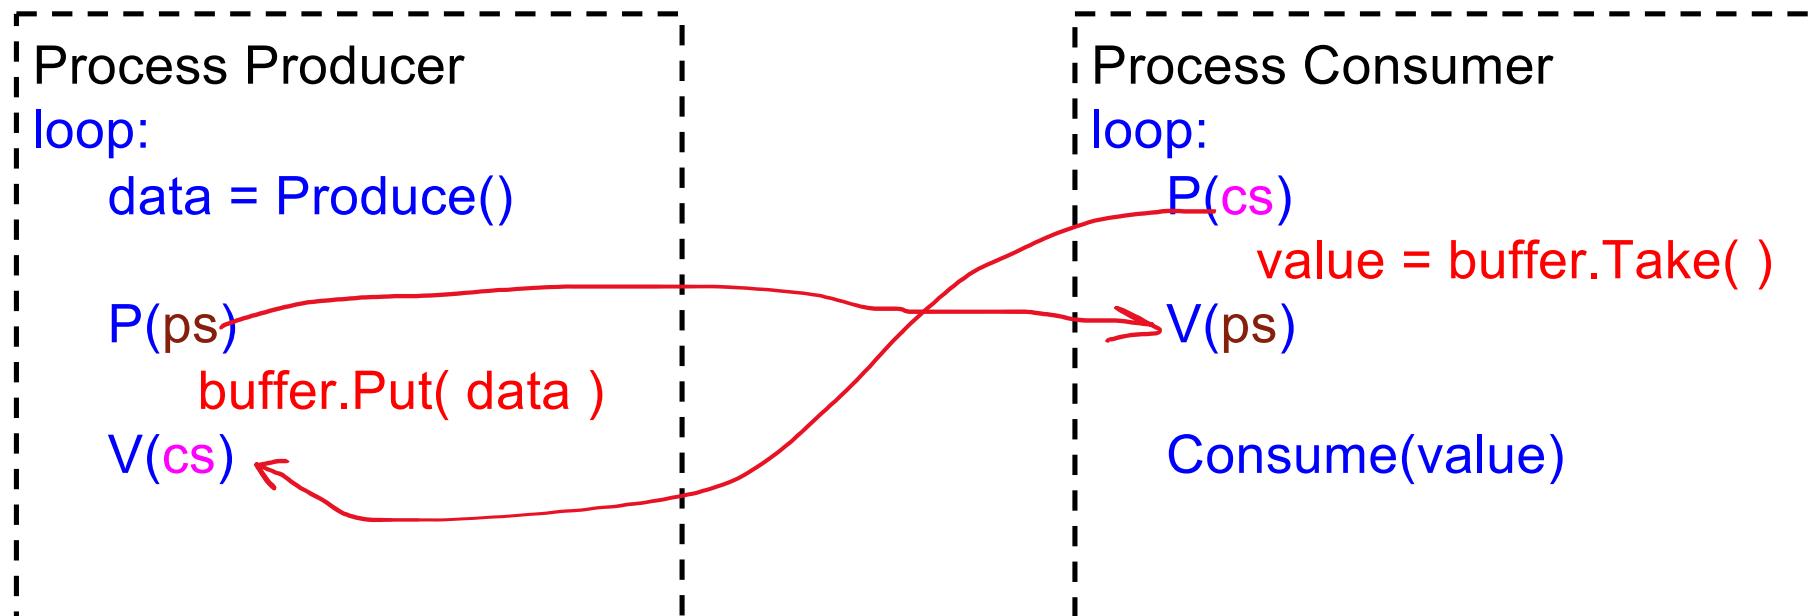
Producer-Consumer

See how semaphores can help us ...

- Assume a shared buffer with only one space
- Define two semaphores: ps and cs
- Initialize $ps=1$ and $cs=0$

Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...



Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...

Process Producer

```
loop:
    data = Produce()

    P(ps)
    buffer.Put( data )
    V(cs)
```

Scenarios:

Start: buffer is empty,
if the consumer is faster, it
has to wait in P(cs). Why?

Process Consumer

```
loop:
    P(cs)
    value = buffer.Take( )
    V(ps)

    Consume(value)
```

Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...

Process Producer

```
loop:  
    data = Produce()  
  
    P(ps)  
    buffer.Put( data )  
    V(cs)
```

Scenarios:

buffer is empty,
if the consumer is faster, it
has to wait in P(cs). Why?

There will be a point
where a producer will
produce data and ...

Process Consumer

```
loop:  
    P(cs)  
    value = buffer.Take( )  
    V(ps)  
    Consume(value)
```

Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...

Process Producer

```
loop:
    data = Produce()

    P(ps)
    buffer.Put( data )
    V(cs)
```

Scenarios:

buffer is empty,
if the consumer is faster, it
has to wait in P(cs). Why?

... it will be the first one to
enter the critical section.

Why?

Process Consumer

```
loop:
    P(cs)
    value = buffer.Take( )
    V(ps)
    Consume(value)
```

Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...

Process Producer

```
loop:  
    data = Produce()  
  
    P(ps)  
    buffer.Put( data )  
    V(cs)
```

Scenarios:

buffer is empty,
if the consumer is faster, it
has to wait in P(cs). Why?

... when there is one
producer in the critical
section, other producers
and all the consumers
have to wait. Why?

Process Consumer

```
loop:  
    P(cs)  
    value = buffer.Take( )  
    V(ps)  
    Consume(value)
```

Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...

Process Producer

```
loop:  
    data = Produce()  
  
    P(ps)  
    buffer.Put( data )  
    V(cs)
```

... when there is one producer in the critical section, other producers and all the consumers have to wait. **Why?**
The producer finishes Put(). It executes V(cs). The waiting consumer will enter the critical section.
Why?

Process Consumer

```
loop:  
    P(cs)  
    value = buffer.Take( )  
    V(ps)  
  
    Consume(value)
```

Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...

Process Producer

```
loop:  
    data = Produce()  
  
    P(ps)  
    buffer.Put( data )  
    V(cs)
```

... when there is one consumer in the critical section, other consumers and all the producers have to wait. **Why?**

The consumer finishes Take(). It executes V(ps). The waiting producer can enter the critical section. **Why?**

Process Consumer

```
loop:  
    P(cs)  
    value = buffer.Take( )  
    V(ps)  
  
    Consume(value)
```

Producer-Consumer

Assume a shared buffer with only one space, we have two semaphores: ps and cs, Initializations are ps=1 and cs=0 ...

Process Producer

```
loop:
    data = Produce()

    P(ps)
    buffer.Put( data )
    V(cs)
```

Interested to implement?
Check the exercises ...

Process Consumer

```
loop:
    P(cs)
    value = buffer.Take( )
    V(ps)

    Consume(value)
```


Producer-Consumer: Extended

We can extend the definition of semaphores to include all positive integers. These semaphores are called counting semaphores.

- Assume a shared buffer with n empty space.
- Define two semaphores: ps and cs .
- Initialize $ps=n$ and $cs=0$.
- ps counts the number of free slots in the buffer, and cs counts the number of data items stored in the buffer.

Producer-Consumer: Extended

- In this example, ps counts the number of free slots in the buffer and cs counts the number of data items inserted into the buffer.

Process Producer

```
loop:
    data = Produce()

    P(ps)
    P(lock)
    buffer.Put( data )
    V(lock)
    V(cs)
```

Here we need another semaphore to provide mutual exclusion in accessing the buffer.

What should be the initial value of the lock?

Process Consumer

```
loop:
    P(cs)
    P(lock)
    value = buffer.Take( )
    V(lock)
    V(ps)

    Consume(value)
```

Part One!

- In this part, we will discuss semaphores and how they can be used to avoid race conditions. Try to find answers to the following questions:
 1. What operations are allowed on a semaphore?
 2. How can a semaphore guarantee mutual exclusion in a critical section?
 3. Do solutions for race conditions using semaphores avoid starvation?
 4. How are counting semaphores used in message-passing systems?
 5. How do you compare semaphore with mutex?

Part Two!


Part Two!

- In this part, we will discuss Deadlocks. Try to find the answer to the following questions:
 1. What are the required conditions to have a deadlock?
 2. What kind of resources do not cause any deadlock?
 3. What mistake in the producer-consumer example can cause a deadlock?
 4. May we have a deadlock if we use a mutex instead of a semaphore in the producer-consumer example?
 5. To prevent deadlocks, the resources should be allocated in a specific order. What order can we use in allocating forks to philosophers?

Deadlocks


Deadlocks

- When each process holds a resource needed by the other process, each process relies on the other process to make a progress.
- This situation when no progress is made is called a **deadlock** condition.



Example 1

- What are the resources here?
- Who is waiting for what?



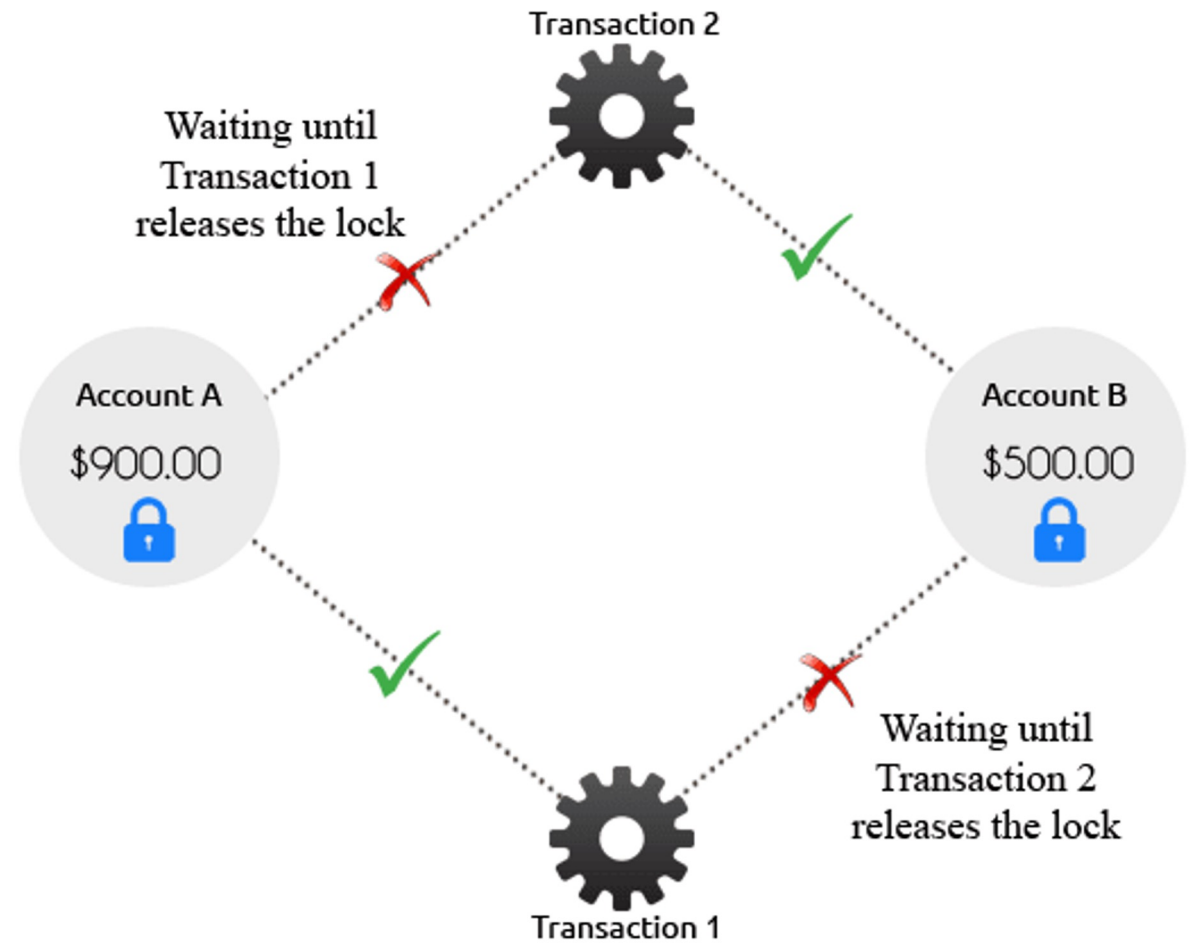
Example 2

- What are the resources here?
- Who is waiting for what?



Example 3

- What are the resources here
- Who is waiting for what?



Conditions to Have a Deadlock

- The four deadlock conditions are:
 - **Exclusive** resource access; **no simultaneous sharing** (only sequentially reusable).
 - **Incremental allocation**: processes do not request all resources together (**hold and wait**).
 - **Circular waiting** for resources.
 - No resource **preemption**.
- These are necessary conditions; if any of them are missing, a deadlock can't occur.

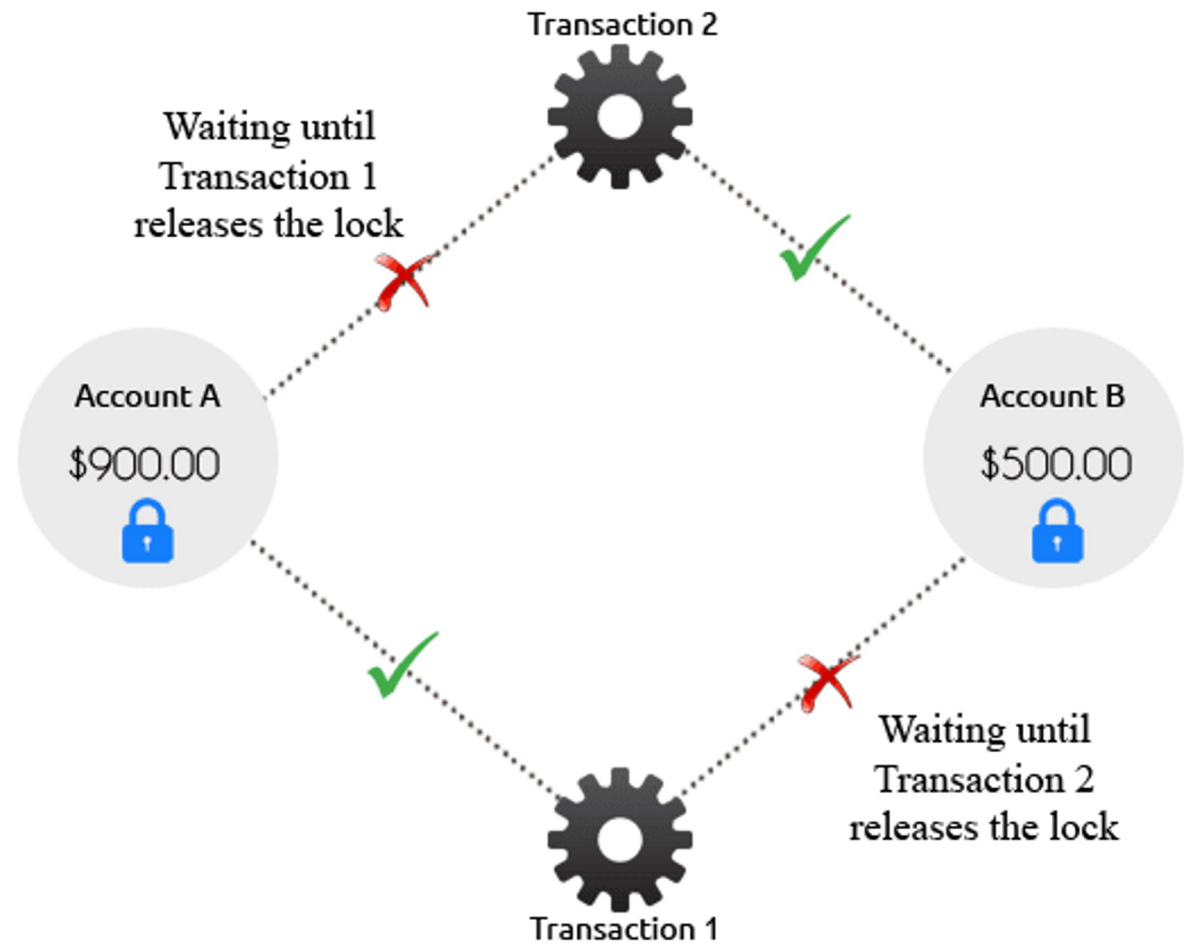
Example 2 (Revised)

- Each process holds a resource and is waiting for the next resource.
- Before allocating resources, the manager should make sure that a circular waiting will not be created. **How?**




Discuss

- How can we avoid deadlocks here?



Deadlock with Semaphores


- A semaphore may be viewed as a resource counter.



- For instance, a semaphore can be used to represent printers.
- If we have two printers, the initial value of the semaphore will be 2.

Deadlock with Semaphores

- Assume processes P and Q are going to use resources R1 and R2.
- P requests the resources in R2, R1 order
- Q requests the resources in R1, R2 order



Discuss



Semaphore $S1=1$, $S2=1$

Process P

P($S1$)

P($S2$)

Run_critical_section_code()

V($S2$)

V($S1$)

Process Q

P($S2$)

P($S1$)

Run_critical_section_code()

V($S1$)

V($S2$)

It is likely to have a deadlock running the above example. Why we are not sure if a deadlock will happen or not?

Mutex and Deadlock

- Deadlocks can happen with Mutex (or Lock) as well
- A Mutex is a lock which can have one of the two states (locked/unlocked)
- Only the process that locks a Mutex can open it.
- A process trying to lock an already locked Mutex is **blocked**

Mutex and Deadlock

- Assume Mutex M is used to enforce exclusive access to a resource



- Process P locks Mutex M



- Then P tries to lock it again.



- But because M is already locked, P will wait until M is unlocked



Discuss



- How long will process P wait? Why?
- What solutions do you suggest?

Deadlock Avoidance

- If any one of the four required deadlock conditions is violated, the deadlock will not happen.
- In practice, we cannot change the type of resource.
- The only condition that the resource manager can control is allocating in a way that circular hold-and-wait queues are not created.

Dining Philosopher

Dining Philosophers is another classical problem which studies Deadlocks ...

- When can a deadlock happen?
- Are you interested in implementing it? Then ..



Extra video → <https://www.youtube.com/watch?v=MYgmmJJfBg>

Part Two!

- In this part, we will discuss Deadlocks. Try to find the answer to the following questions:
 1. What are the required conditions to have a deadlock?
 2. What kind of resources do not cause any deadlock?
 3. What mistake in the producer-consumer example can cause a deadlock?
 4. May we have a deadlock if we use a mutex instead of a semaphore in the producer-consumer example?
 5. To prevent deadlocks, the resources should be allocated in a specific order. What order can we use in allocating forks to philosophers?



Time to Apply

Week 5: Exercisez