

Concurrency

Week 2

Processes

Contents

- Processes
 - Definition
 - Creation
 - Termination
- Scheduling and Context Switching
 - Context Switching
 - Long-Term and Short-Term Scheduling
- Inter-process communication

Learning Outcomes of this Lesson

- At the end of this lesson, you will be able to:
 - Describe the concept of a process, its creation, its changing state, and its termination.
 - Understand multitasking and context-switching
 - Understand process types in terms of inter-process communications

Part One: Processes


- In this part, you will be learning the concept of a process and its life cycle. Try to find answers to the following questions:
 1. What is a process?
 2. What does a process need to execute (resources needed)
 3. How is a process created?
 4. What is a process state, and how does a process state change?
 5. How does a process terminate?

Process Concept: What is a Process?

- An operating system executes a variety of programs:
 - Text editors
 - Data processing
 - Web applications
 - Interactive games
- Process – a **program** in **execution**; process execution must progress in a sequential fashion

How a Program is Put in Execution?

- To put a program into **execution** we should:
 - Copy it (at least partially!) into the main memory
 - Set the program counter (pc) to point to the first instruction
 - Initialize stack
 - Reset flags
 - The process is ready, start running ...



What is the difference between a program and a process?

- Process – a **program** in **execution**; process execution must progress sequentially (as from the slides before...). It is an active entity born from a program.
- A program can have multiple processes (initiates/owns)
- Multiple processes can be part of a program
- A process performs a task

What Resources are Used by a Process?

- A process has its own:
 - Code, data, and stack
 - Usually (but not always, when not?) has its own address space
 - Program State
 - CPU registers
 - The program counter (current location in the code)
 - Stack pointer
- Only one process can be running in the CPU at any given time!

When is a Process Created?

- Processes can be created in two ways
 - **System initialization**: some processes are created when the OS starts up
 - By executing a process creation **system call**: another process **explicitly** asks for a new process
- System calls can come from:
 - ***A user “request”*** to create a new process (system call executed from a user shell, for instance when a user clicks on an icon on the desktop)
 - ***already running processes***
 - User programs
 - System daemons*

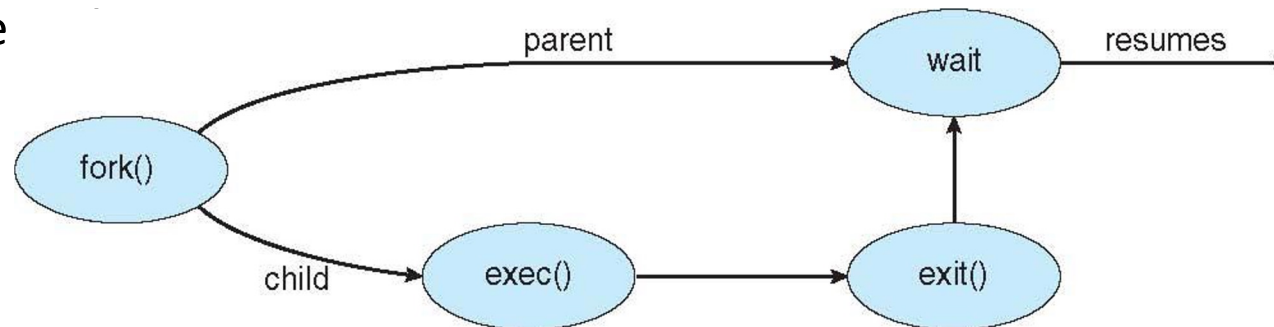
*System daemons: process that runs in background

Process Creation Options

- Resource-sharing options
 - Parent and children share all resources
 - Children share a subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation Options (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space



When do Processes End?

- Conditions that terminate processes can be
 - Voluntary
 - Involuntary
- Voluntary
 - Normal exit: when the process has no more instruction to execute.
 - Error exit
- Involuntary
 - Fatal error
 - Killed by another process

Process Termination

- The process executes the last statement and then asks the operating system to delete it using the `exit()` system call.
 - Returns `status` data from the child to the parent (via `wait()`)
 - Process `resources` are deallocated by the operating system
- The parent may terminate the execution of the children's processes using the `abort()` system call.
 - Some reasons for doing so:
 - A child has exceeded the allocated resources
 - Task assigned to a child is no longer required
 - The parent is exiting, and the operating system does not allow a child to continue if its parent terminates

Process Termination

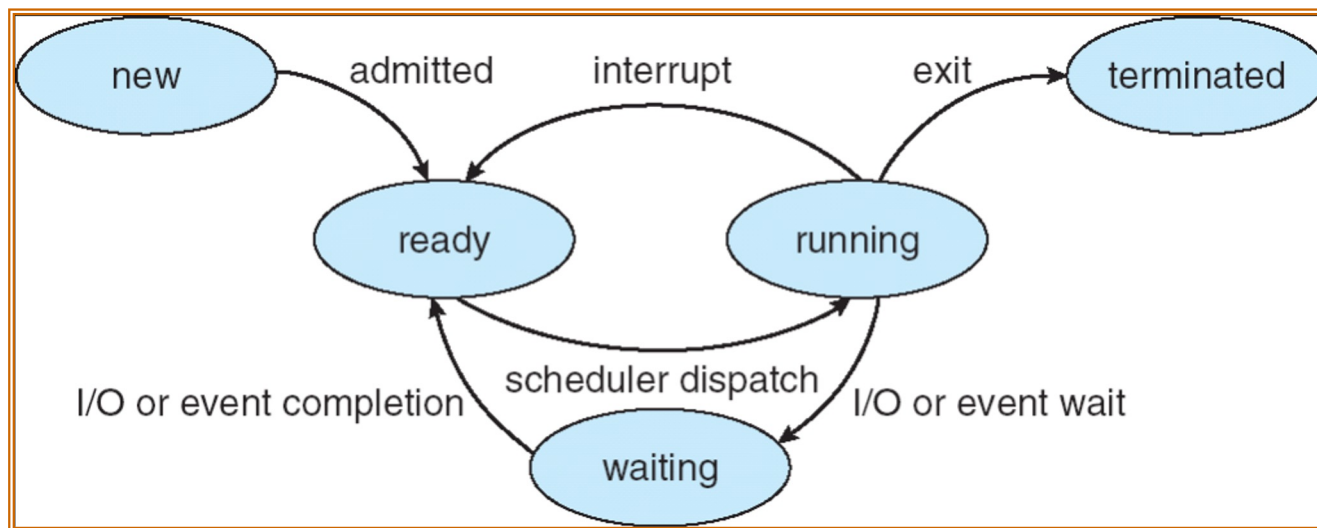
- Some operating systems do not allow a child to exist if its parent has terminated. If a process *terminates*, then all its children must also be terminated.
 - **cascading termination**. All children, grandchildren, etc. are terminated.
 - The termination is initiated by the **operating system**.
- The parent process may wait for the termination of a child process by using the **wait()** system call.
 - The call returns status information and the PID of the terminated process
PID = wait(&status);
- If no parent is waiting (did not invoke **wait()**, *yet*), process is a **zombie (*NIX)**
- If **parent** is **terminated** without invoking **wait()**, process is an **orphan (*NIX)**

Process hierarchies

- Parent creates a child process
 - Child processes can create their own children
- Forms a hierarchy
 - UNIX calls this a “process group”
 - If a process exits, its children are “inherited” by the exiting process’s parent
- Windows has no concept of process hierarchy (in its base definition)
 - All processes are created equal
 - we can specify ***now*** how the inheritance kind with code (Win10)

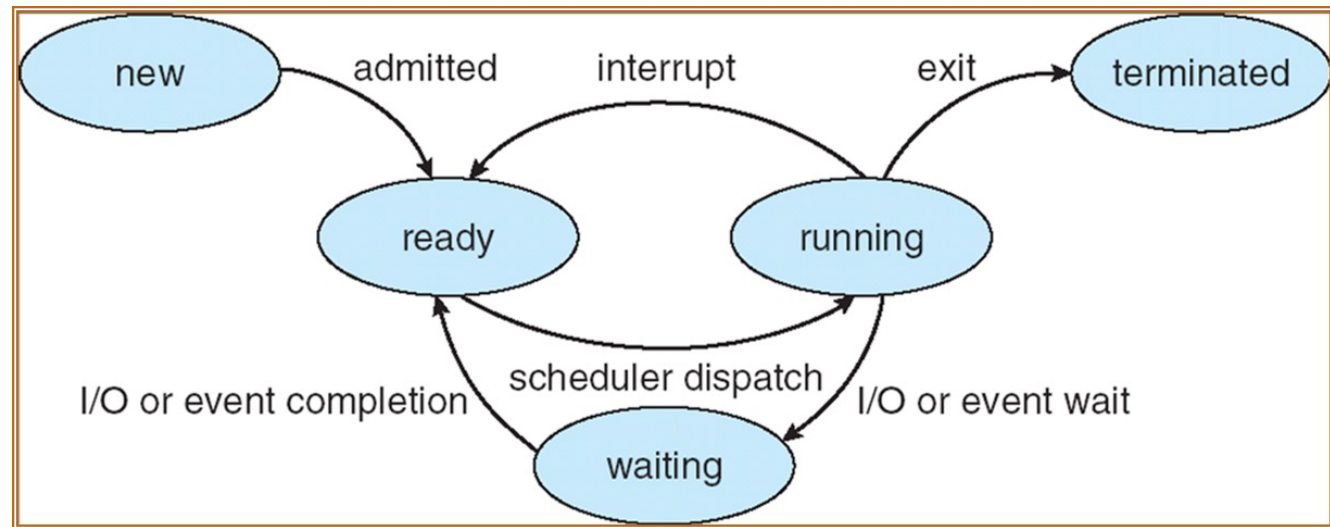
Process State

- As a process executes, it changes its **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a CPU
 - **terminated**: The process has finished execution



Process State Transition

- Possible transitions between states are
 - 1 - Process enters **ready** queue
 - 2 - Scheduler picks this process (change to **running** state)
 - 3 - Scheduler picks a different process (go back to **ready** state)
 - 4 - Process **waits** for event (such as I/O)
 - 5 - Event occurs (go to **waiting** state)
 - 6 - Process exits (**terminated**)
 - 7 - Process ended by another process (**terminated**)



Part One: Processes

- Now you can give answers to the following questions:
 1. What is a process?
 2. What does a process need to execute (resources needed)
 3. How is a process created?
 4. What is a process state, and how does a process state change?
 5. How does a process terminate?

Part Two!

Part Two: Scheduling and Context Switching


- In this part, you will learn how processes share CPU time. Pay attention to find the answers to the following questions:
 1. What is meant by CPU time sharing?
 2. Why do we need time-sharing?
 3. What is context-switching?
 4. What is the context of a process and how is it stored?
 5. What is process scheduling? What is a simple process scheduling model?

Switching from Process to Process


- Running multiple processes at the same time requires sharing CPU time (Assuming a single CPU/core)
- Stopping a process and allocating CPU to another process is called **context switching**
- Context switching is used for:
 - Maximizing **CPU utilization** by running a ready process when the current process is blocked (waiting for I/O for instance)
 - Creating the **illusion of a multi-processor** system on single-processor systems by running each process for a short period named a **quantum**.
- Context-switch time is **overhead**; the system does no useful work while switching context

The Process Model: Multiprogramming


- Multiprogramming of 4 programs
- Conceptual model
 - 4 independent processes
 - Processes run sequentially
- Only one program is active at any instant!
 - That instant can be very short...



Single PC
(CPU's point of view)

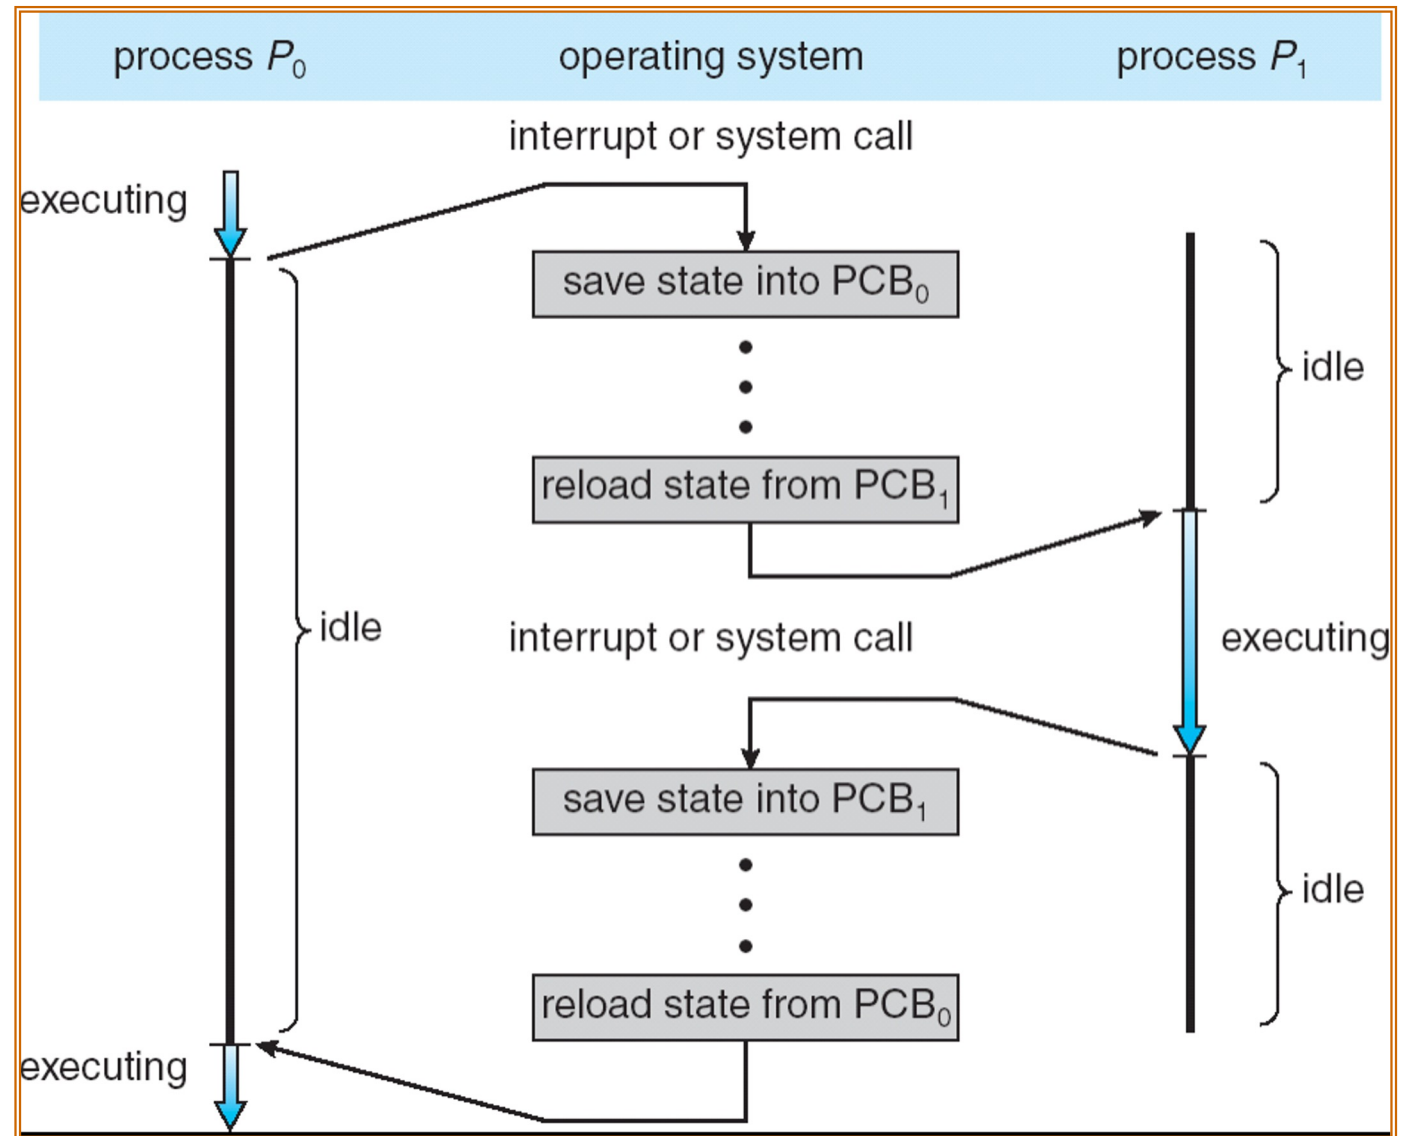


Multiple PCs
(process point of view)



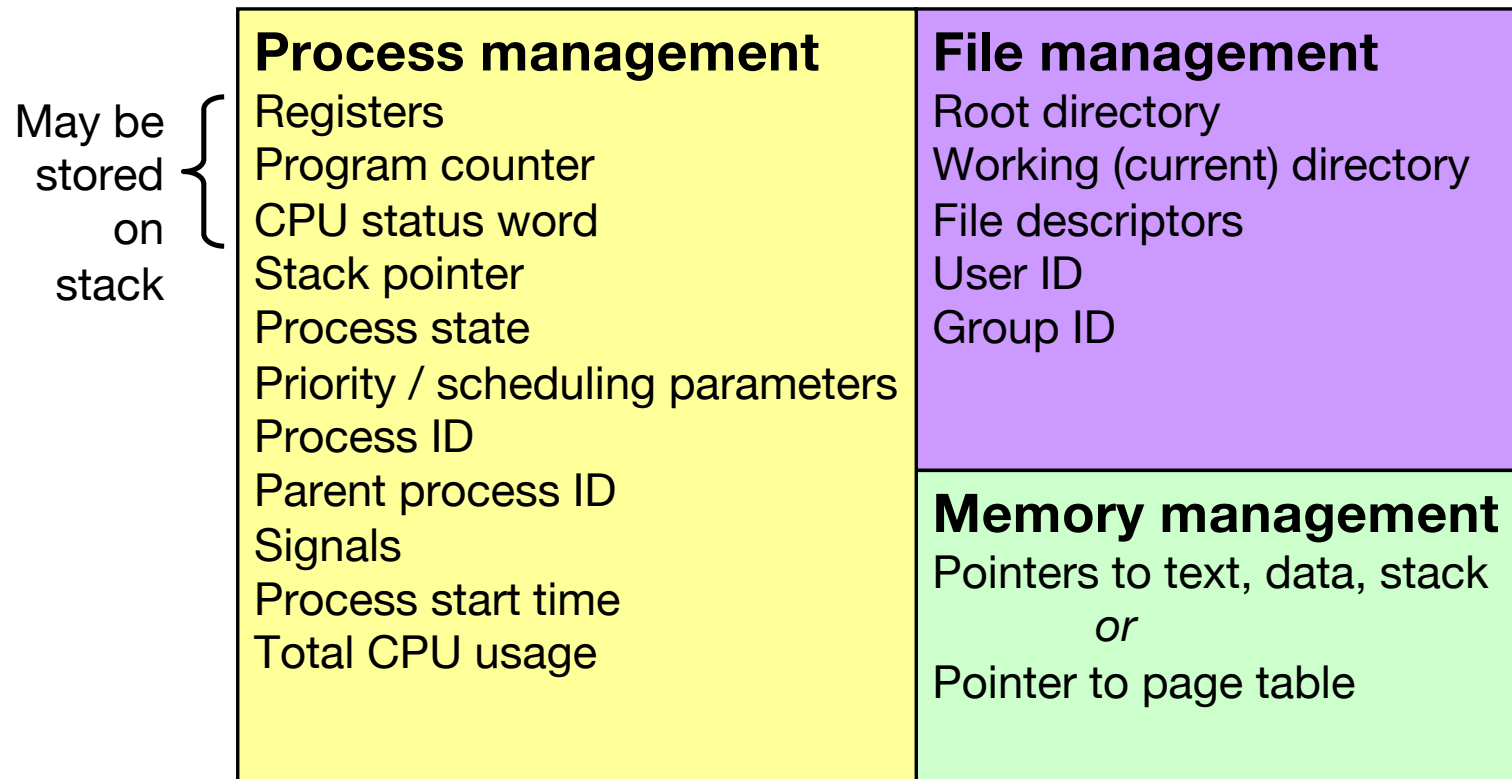
Context Switching

PCB: process control block (see op4 year 1)



What is a Process Control Block?

All information about a process is stored in a data structure named Process Control Block (PCB)



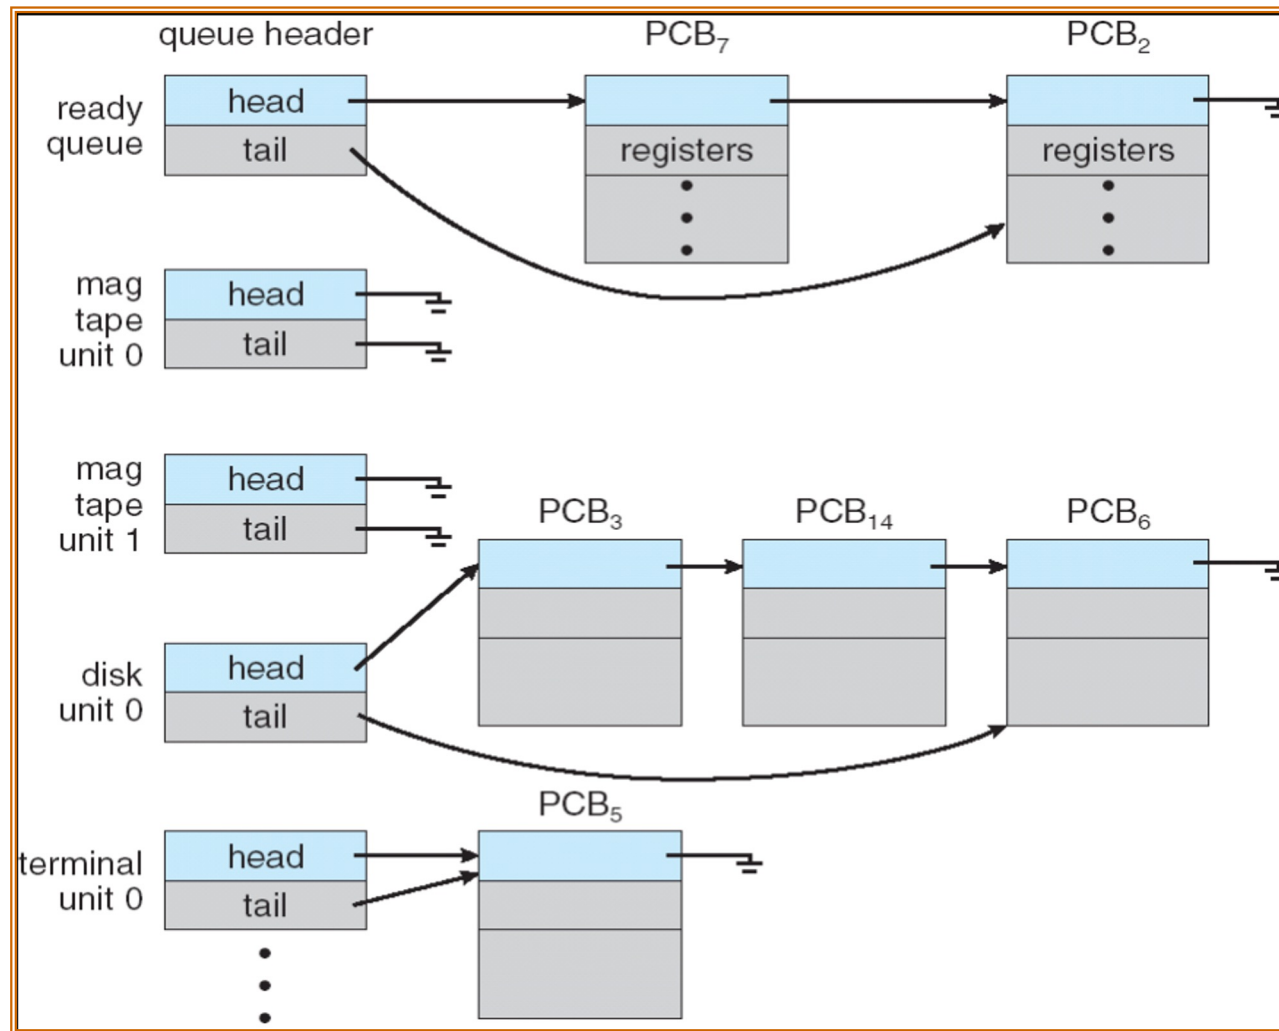
Process Scheduling

- Context switching requires choosing a process to allocate CPU when the CPU is withheld from the current process.
- **Process scheduling** refers to the methods used for selecting the next process to run (see Operating System course for the details)
- To apply process scheduling algorithms, multiple **queues** are created

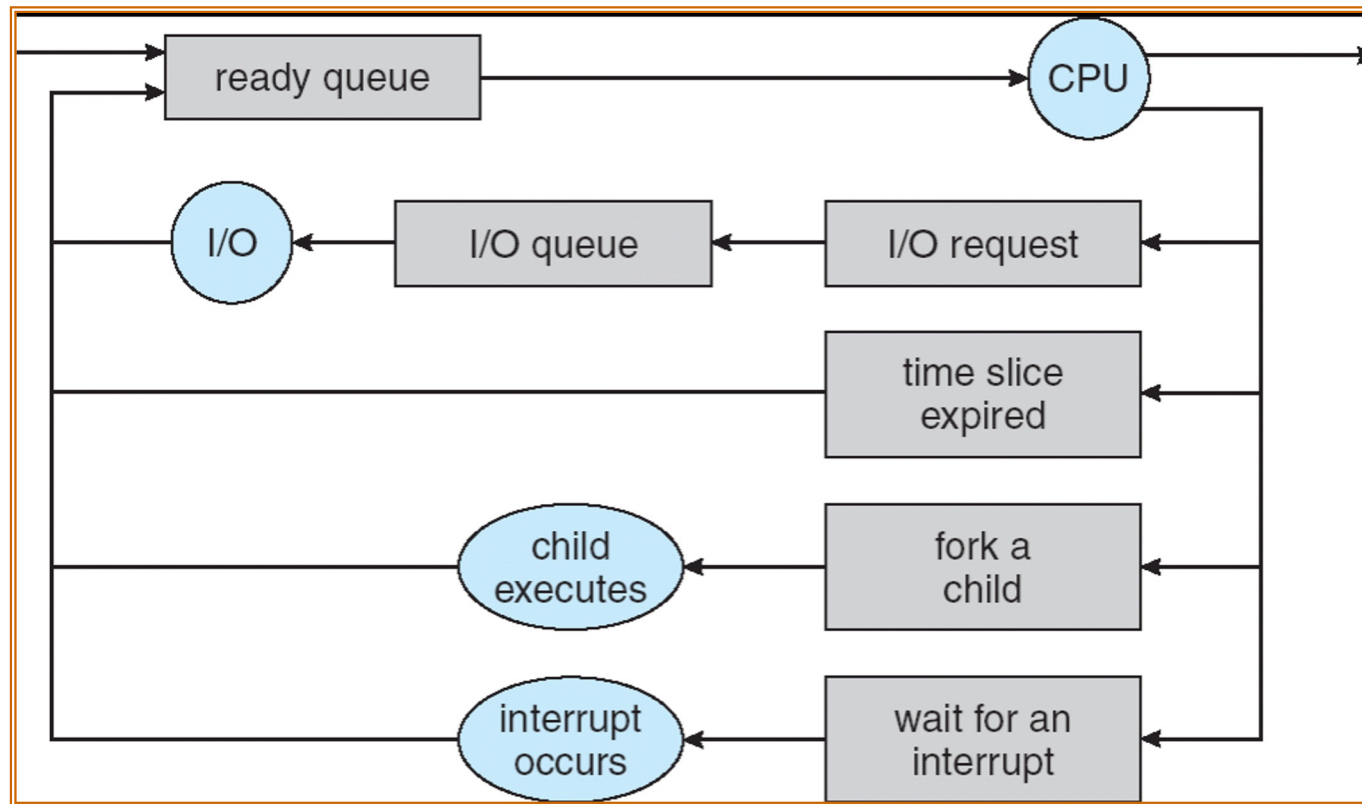
Process Scheduling Queues

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to be executed
- **Device queues** – a set of processes waiting for an I/O device
- Processes migrate among the various queues

Ready Queue and Various I/O Device Queues



Simplified Model of Process Scheduling



Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the *ready queue from the job queue*
 - The long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
 - The long-term scheduler controls the *degree of multiprogramming*
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
 - The short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts

Part Two: Scheduling and Context Switching

- Now you can give answers for the following questions:
 1. What is meant by CPU time sharing?
 2. Why do we need time-sharing?
 3. What is context-switching?
 4. What is the context of a process and how is it stored?
 5. What is process scheduling? What is a simple process scheduling model?

Part Three!

Part Three: Inter-process communications

- In this part, you will learn about communication between processes. Try to find answers to the following questions:
 1. What is meant by communication between processes?
 2. Does every process have to communicate with other processes?
 3. What are the main methods for communication between processes?

Process Classes (kinds)

- **The independent** process cannot affect or be affected by the execution of another process
- **The cooperating** process can affect or be affected by the execution of another process
- Advantages/objectives of process cooperation
 - Information Sharing
 - Computational speed-up (given the right conditions)
 - Modularity (ease of maintenance and abstractions of tasks)
 - Convenience

Example: Cooperating Processes

- Producer-Consumer Problem

- A process named *producer* generates data items. A second process named *consumer* utilizes the data
- *producer* process shares the information with the *consumer* process using a shared buffer which can be either a(n):
 - *Unbounded buffer* places no practical limit on the size of the buffer
 - *bounded buffer* assumes that there is a fixed buffer size

Inter-Process Communication (IPC)

- The mechanism for processes to communicate and synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)
- If *P* and *Q* wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via *send/receive*
- Implementation of a communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

Facilitating Inter-Process Communication

- If the cooperating processes share information for computation speed-up or modularity, they can be part of the same process.
- Being part of the same process makes it possible to
 - Share the code
 - Share the data
 - Share the resources
- However, each one should have its own
 - Stack
 - Registers
 - Program counter
- Each sub-part of a process is named a **thread** (next week)

Interprocess Communication

- Processes within a system may be *independent* or *cooperating*
- The cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **inter-process communication (IPC)**
- Two models of IPC
 - **Shared memory**
 - **Message passing**

Part Three: Inter-process communications

- Now, you can give answers to the following questions:
 1. What is meant by communication between processes?
 2. Does every process have to communicate with other processes?
 3. What are the main methods for communication between processes?

Summary

- A **process** is a program in execution
- A process is **created** at system initialization, or through a system call issued by another process
- Processes may **terminate** voluntarily or be killed by another process
- Processes may take different **states** during execution
- The operating system revokes the CPU from the running process and grants it to another process, a procedure called **context switching**
- Operating systems use different **scheduling algorithms** to decide which process should run next

Quiz Time

- Please, answer the quiz questions for week two.
- You have 10 minutes.
- We will discuss